

Compiler 2017

Linux x86-64 Assembly in NASM

Lequn Chen
March 16, 2017

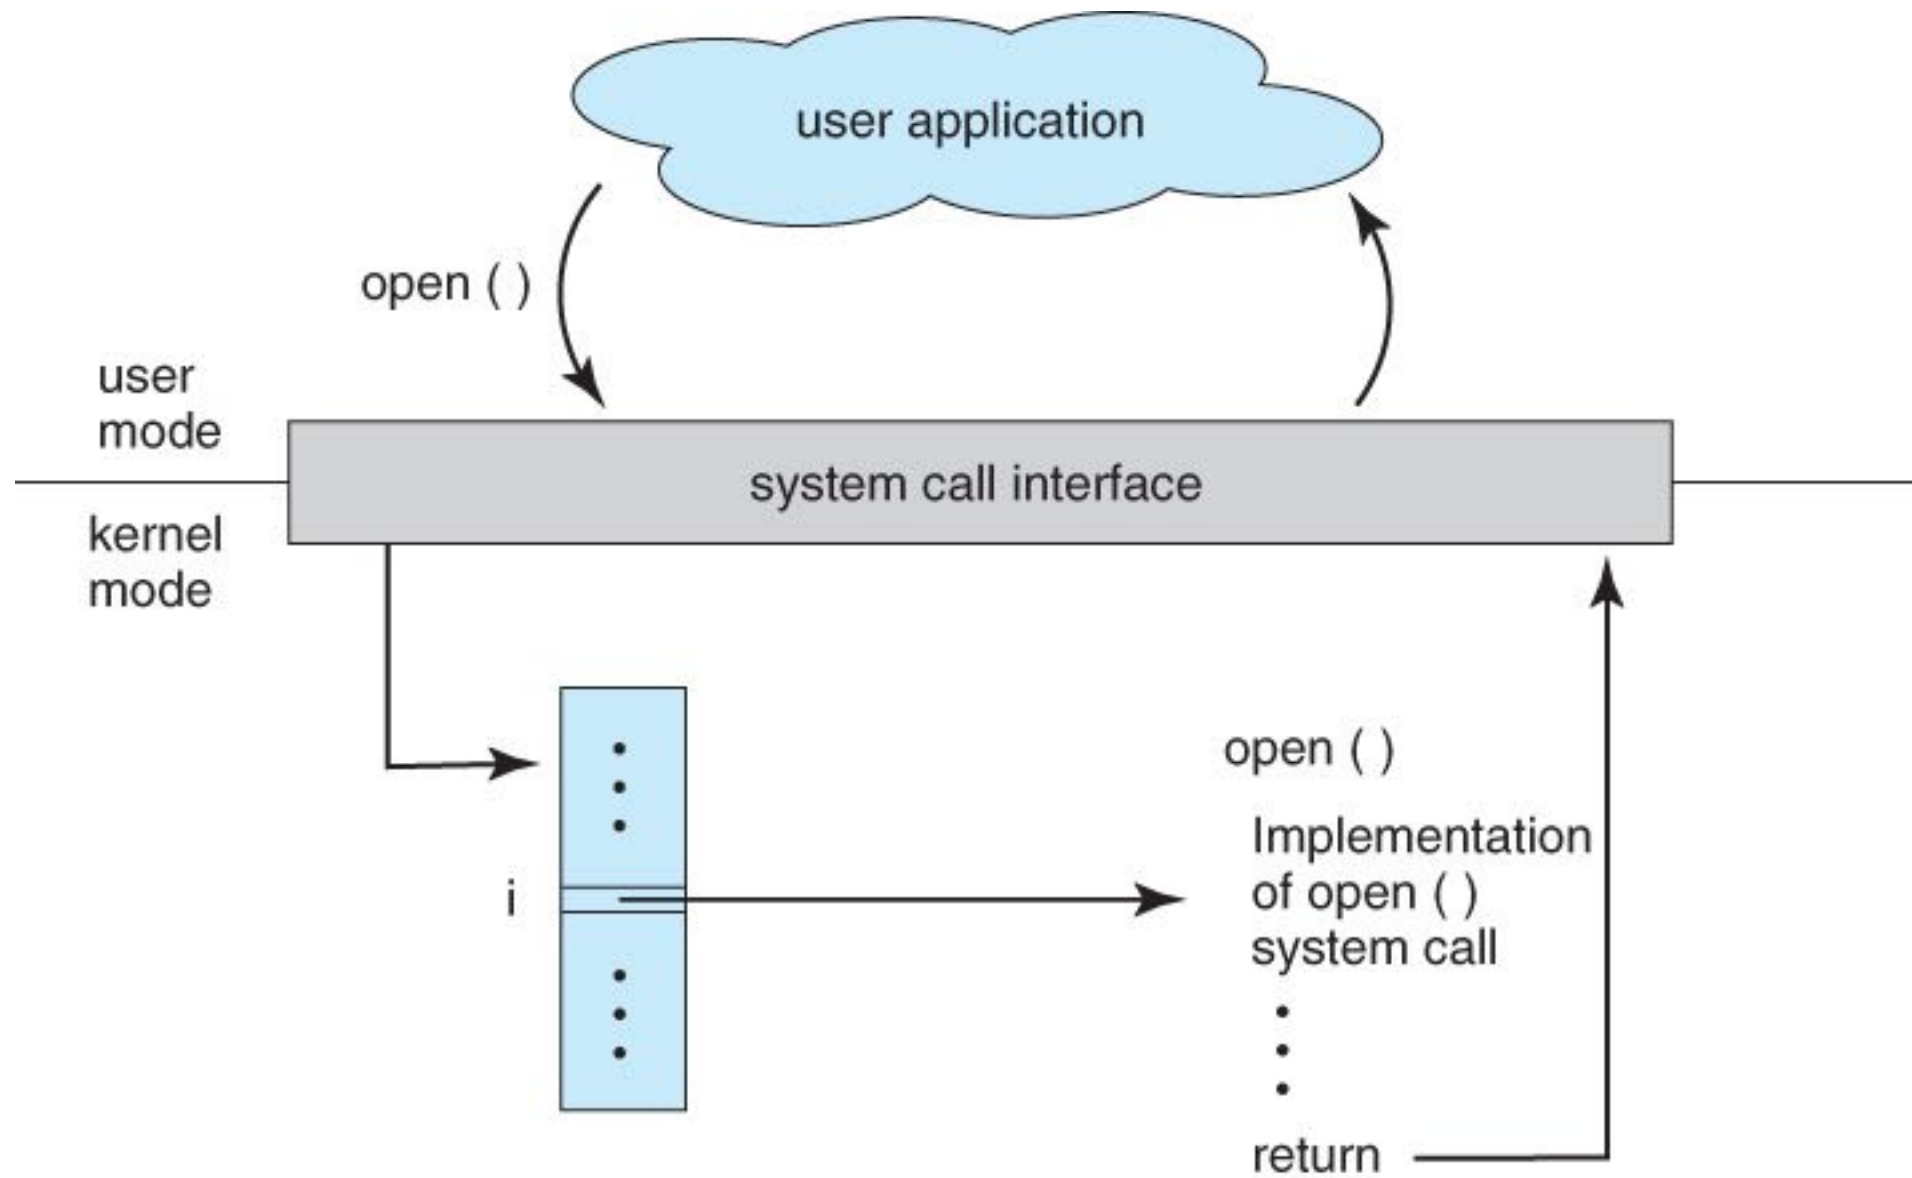
Operating System

- User Space
 - Your program runs in user space
 - Limited capability (simply put: an advanced calculator)
 - Calls some “magic function” to obtain super power! (e.g. read, write, ...)
- Kernel Space
 - Dirty work: Deal with hardwares (e.g. raw I/O, interrupt, ...)
 - Supervise: Kill user programs when unauthorized action
 - Service: Provide easy-to-use APIs
 - The most low level API: system call

C Library

- libc: some most basic high level operations
 - scanf / printf
 - memset / memcpy
 - strlen / strcmp
 - ...
- in user space
- libc calls syscall for you, so you don't feel their existence

Operating System



```
#include <stdio.h>
```

```
int main(void) {  
    char name[10];  
    scanf("%s", name);  
    printf("Hello, %s\n", name);  
    return 0;  
}
```

Libc

```
#include <unistd.h>
```

```
int main(void) {  
    char name[10];  
    ssize_t n = read(0, name, 10);  
    write(1, "Hello, ", 8);  
    write(1, name, n);  
    write(1, "\n", 2);  
    return 0;  
}
```

Linux API

```
#include <unistd.h>
```

```
#include <sys/syscall.h>
```

```
int main(void) {  
    char name[10];  
    long n = syscall(3, 0, name, 10);  
    syscall(4, 1, "Hello, ", 8);  
    syscall(4, 1, name, n);  
    syscall(4, "\n", 2);  
    return 0;  
}
```

Linux System Call

Hello World in NASM

```
; -----
; Writes "Hello, World" to the console using only system calls. Runs on 64-bit Linux only.
; To assemble and run:
;
;     nasm -felf64 hello.asm && ld hello.o && ./a.out
; -----
```

```
global _start

section .text

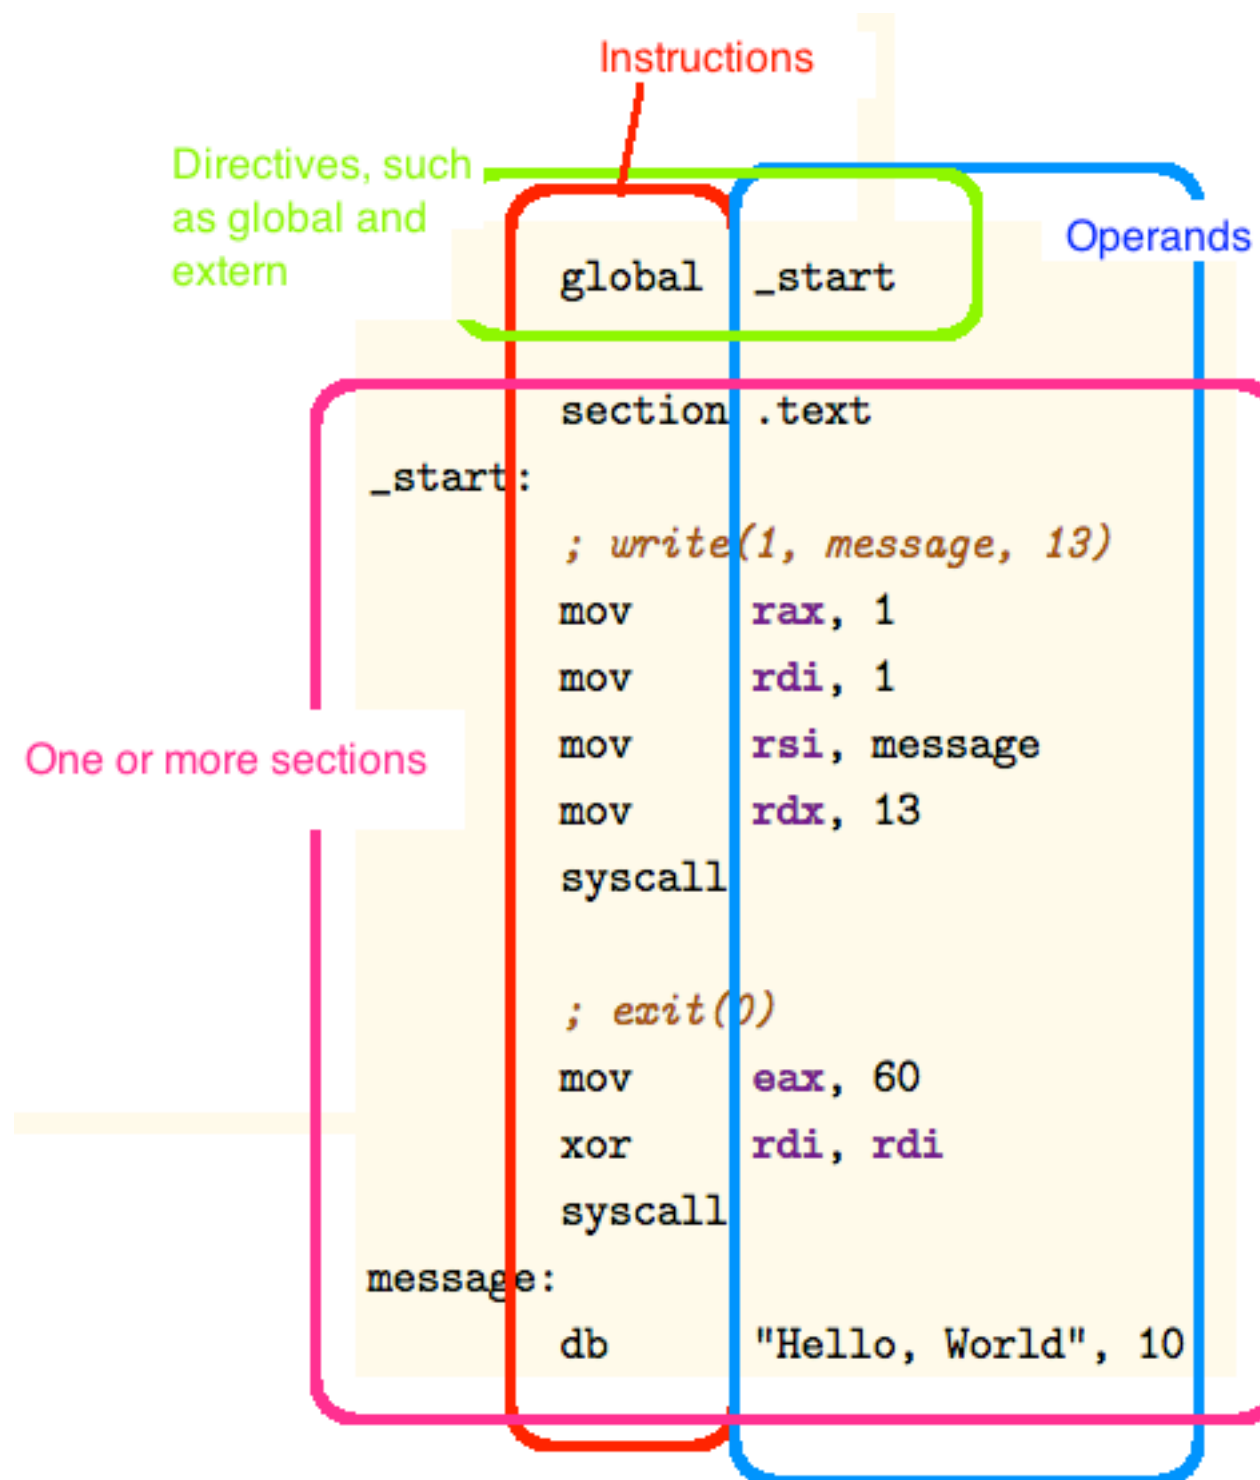
_start:
    ; write(1, message, 13)
    mov     rax, 1           ; system call 1 is write
    mov     rdi, 1           ; file handle 1 is stdout
    mov     rsi, message     ; address of string to output
    mov     rdx, 13          ; number of bytes
    syscall                 ; invoke operating system to do the write

    ; exit(0)
    mov     eax, 60          ; system call 60 is exit
    xor     rdi, rdi         ; exit code 0
    syscall                 ; invoke operating system to exit

message:
    db      "Hello, World", 10 ; note the newline at the end
```

<pre>mov x, y: x := y xor x, y: x := x xor y syscall : invoke system call</pre>

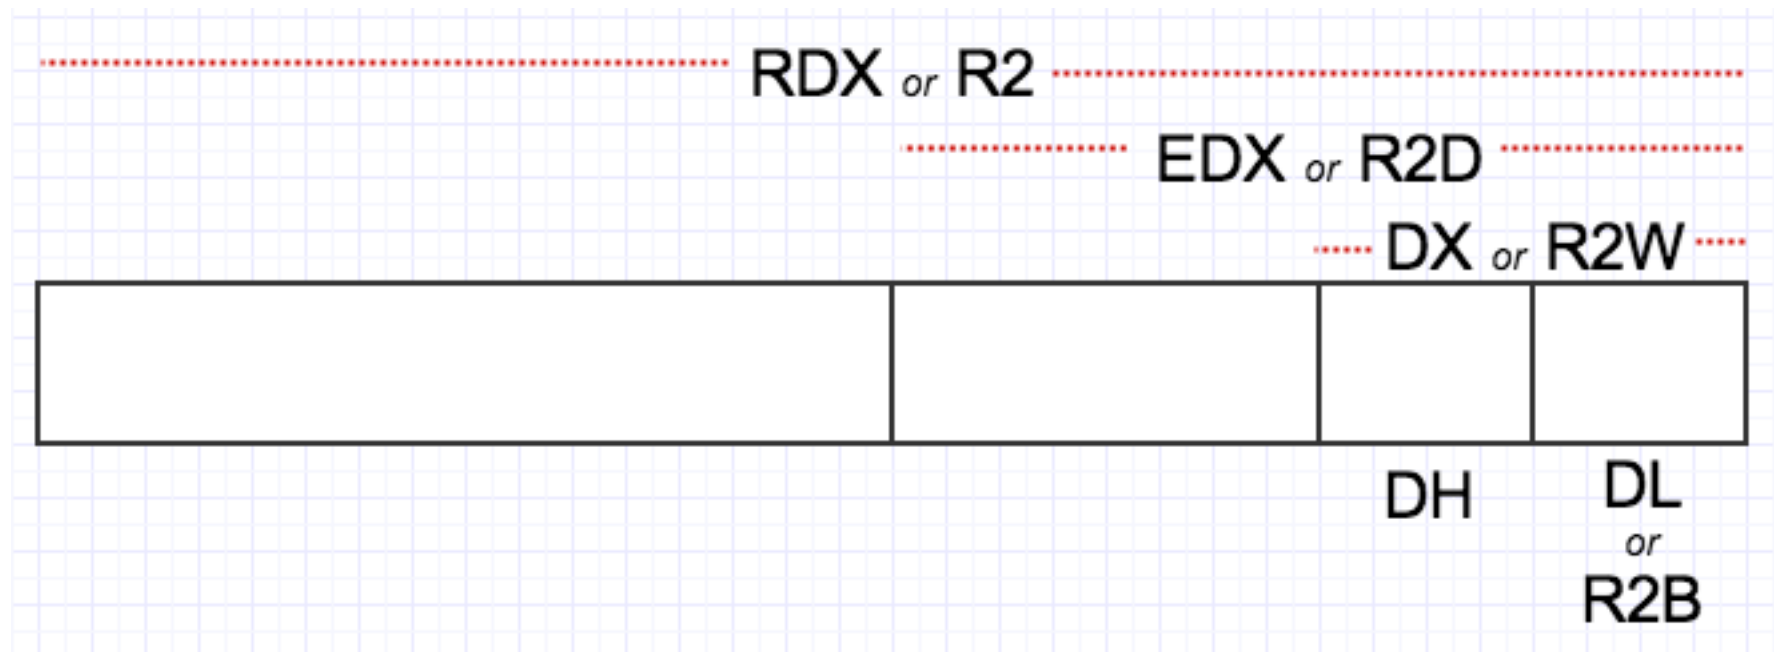
Structure of a NASM Program



Register Operands

- **64-bit**
 - RAX RCX RDX RBX RSP RBP RSI RDI R8 R9 R10 R11 R12 R13 R14 R15
R0 R1 R2 R3 R4 R5 R6 R7
- **32-bit**
 - EAX ECX EDX EBX ESP EBP ESI EDI R8D R9D R10D R11D R12D R13D R14D R15D
R0D R1D R2D R3D R4D R5D R6D R7D
- **16-bit**
 - AX CX DX BX SP BP SI DI R8W R9W R10W R11W R12W R13W R14W R15W
R0W R1W R2W R3W R4W R5W R6W R7W
- **8-bit (low)**
 - AL CL DL BL SPL BPL SIL DIL R8B R9B R10B R11B R12B R13B R14B R15B
R0B R1B R2B R3B R4B R5B R6B R7B
- **8-bit (high)**
 - AH CH DH BH

Register Operands



Memory Operands

base	[number]	
	[reg]	
index	[reg + reg*scale]	scale in {1,2,4,8}
displacement	[reg + number]	
	[reg + reg*scale + number]	

[750]	; displacement only
[rbp]	; base register only
[rcx + rsi*4]	; base + index * scale
[rbp + rdx]	; scale is 1
[rbx - 8]	; displacement is -8
[rax + rdi*8 + 500]	; all four components

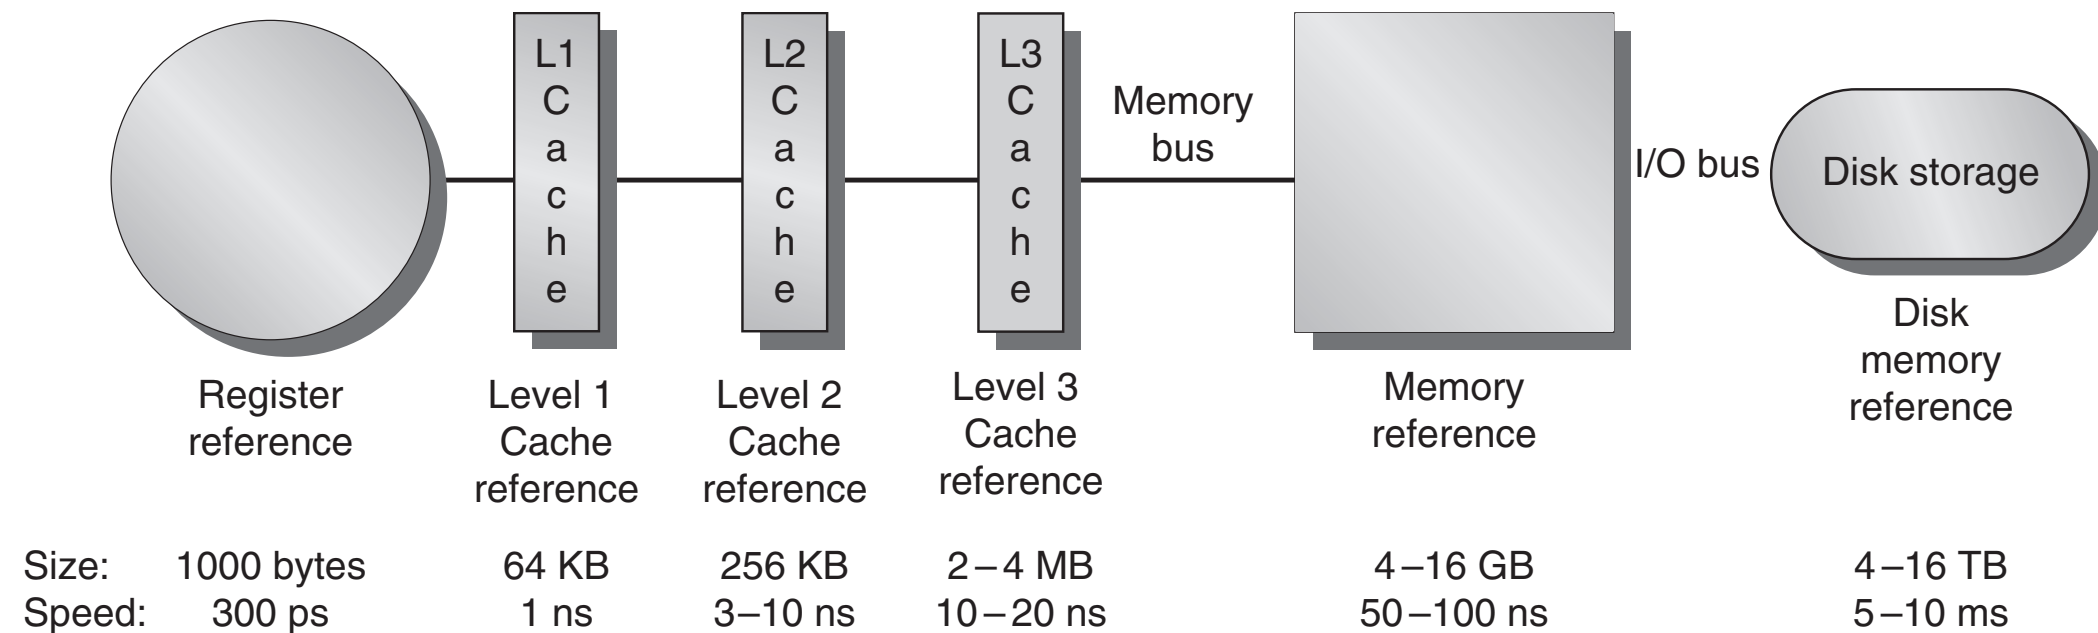
Immediate Operands

```
200          ; decimal
0200         ; still decimal - the leading 0 does not make it octal
0200d        ; explicitly decimal - d suffix
0d200        ; also decimal - 0d prefix
0c8h         ; hex - h suffix, but leading 0 is required because c8h looks like a var
0xc8         ; hex - the classic 0x prefix
0hc8         ; hex - for some reason NASM likes 0h
310q         ; octal - q suffix
0q310        ; octal - 0q prefix
11001000b    ; binary - b suffix
0b1100_1000  ; binary - 0b prefix, and by the way, underscores are allowed
```

Instructions

- Instructions with two memory operands are extremely rare. Most of the basic instructions have only the following forms:
 - `add reg, reg`
 - `add reg, mem`
 - `add reg, imm`
 - `add mem, reg`
 - `add mem, imm`

Memory Hierarchy



(a) Memory hierarchy for server


Defining Data

db	0x55	; just the byte 0x55
db	0x55,0x56,0x57	; three bytes in succession
db	'a',0x55	; character constants are OK
db	'hello',13,10,'\$'	; so are string constants
dw	0x1234	; 0x34 0x12
dw	'a'	; 0x61 0x00 (it's just a number)
dw	'ab'	; 0x61 0x62 (character constant)
dw	'abc'	; 0x61 0x62 0x63 0x00 (string)
dd	0x12345678	; 0x78 0x56 0x34 0x12
dd	1.234567e20	; floating-point constant
dq	0x123456789abcdef0	; eight byte constant
dq	1.234567e20	; double-precision float
dt	1.234567e20	; extended-precision float

Reserving Space

```
buffer:          resb    64          ; reserve 64 bytes
wordvar:         resw    1          ; reserve a word
realarray:       resq   10          ; array of ten reals
```

Data Segments



Using a C Library

```
; -----
; Writes "Hola, mundo" to the console using a C library. Runs on Linux or any other system
; that does not use underscores for symbols in its C library. To assemble and run:
;
;     nasm -felf64 hola.asm && gcc hola.o && ./a.out
; -----
```

```
global main
extern puts

section .text

main:                                ; This is called by the C library startup code
    mov     rdi, message             ; First integer (or pointer) argument in rdi
    call    puts                    ; puts(message)
    ret                               ; Return from main back into C library wrapper

message:
    db      "Hola, mundo", 0        ; Note strings must be terminated with 0 in C
```

call: call a function



Calling Conventions

- From left to right, pass as many parameters as will fit in registers. The order in which registers are allocated, are: `rdi, rsi, rdx, rcx, r8, r9`
- Additional parameters are pushed on the stack, right to left, and are to be removed by the caller after the call.
- After the parameters are pushed, the call instruction is made, so when the called function gets control, the return address is at `[rsp]`, the first memory parameter is at `[rsp+8]`, etc.

Calling Conventions

- The stack pointer `rsp` must be aligned to a 16-byte boundary before making a call.
- The only registers that the called function is required to preserve (the callee-save registers) are: `rbp`, `rbx`, `r12`, `r13`, `r14`, `r15`. All others are free to be changed by the called function.
- Integers are returned in `rax`

Calling Conventions



```
long myfunc(long a, long b, long c, long d,
            long e, long f, long g, long h)
{
    long xx = a * b * c * d * e * f * g * h;
    long yy = a + b + c + d + e + f + g + h;
    long zz = utilfunc(xx, yy, xx % yy);
    return zz + 20;
}
```

Register Usage

Register	Usage	Preserved across function calls
%rax	temporary register; with variable arguments passes information about the number of vector registers used; 1 st return register	No
%rbx	callee-saved register	Yes
%rcx	used to pass 4 th integer argument to functions	No
%rdx	used to pass 3 rd argument to functions; 2 nd return register	No
%rsp	stack pointer	Yes
%rbp	callee-saved register; optionally used as frame pointer	Yes
%rsi	used to pass 2 nd argument to functions	No
%rdi	used to pass 1 st argument to functions	No
%r8	used to pass 5 th argument to functions	No
%r9	used to pass 6 th argument to functions	No
%r10	temporary register, used for passing a function's static chain pointer	No
%r11	temporary register	No
%r12-r14	callee-saved registers	Yes
%r15	callee-saved register; optionally used as GOT base pointer	Yes

Fibonacci Example

```
; -----
; A 64-bit Linux application that writes the first 90 Fibonacci numbers. To
; assemble and run:
;
;     nasm -felf64 fib.asm && gcc fib.o && ./a.out
; -----
```

```
        global main
        extern printf

        section .text
main:
        push    rbx                ; we have to save this since we use it

        mov     ecx, 90            ; ecx will countdown to 0
        xor     rax, rax           ; rax will hold the current number
        xor     rbx, rbx           ; rbx will hold the next number
        inc     rbx               ; rbx is originally 1
```

```
push x    : rsp -= sizeof(x); [rsp] := x;
pop x     : x := [rsp]; rsp += sizeof(x);
jnz label : if Z flag is set, jump to label
call label: push address of next inst; jump to label
ret       : Pop into the instruction pointer
```

Fibonacci Example

print:

```
; We need to call printf, but we are using rax, rbx, and rcx.  printf
; may destroy rax and rcx so we will save these before the call and
; restore them afterwards.
```

```
push    rax                ; caller-save register
push    rcx                ; caller-save register

mov     rdi, format        ; set 1st parameter (format)
mov     rsi, rax           ; set 2nd parameter (current_number)
xor     rax, rax           ; because printf is varargs

; Stack is already aligned because we pushed three 8 byte registers
call    printf            ; printf(format, current_number)

pop     rcx               ; restore caller-save register
pop     rax               ; restore caller-save register

mov     rdx, rax          ; save the current number
mov     rax, rbx          ; next number is now current
add     rbx, rdx          ; get the new next number
dec     ecx               ; count down
jnz     print            ; if not done counting, do some more

pop     rbx               ; restore rbx before returning
ret
```

format:

```
db     "%20ld", 10, 0
```

Recursion Example

```

; -----
; An implementation of the recursive function:
;
; uint64_t factorial(uint64_t n) {
;     return (n <= 1) ? 1 : n * factorial(n-1);
; }
; -----

```

```

global factorial

```

```

section .text

```

```

factorial:

```

cmp	rdi, 1	; n <= 1?
jnbe	L1	; if not, go do a recursive call
mov	rax, 1	; otherwise return 1
ret		
L1:		
push	rdi	; save n on stack (also aligns %rsp!)
dec	rdi	; n-1
call	factorial	; factorial(n-1), result goes in %rax
pop	rdi	; restore n
imul	rax, rdi	; n * factorial(n-1), stored in %rax
ret		

Heap?

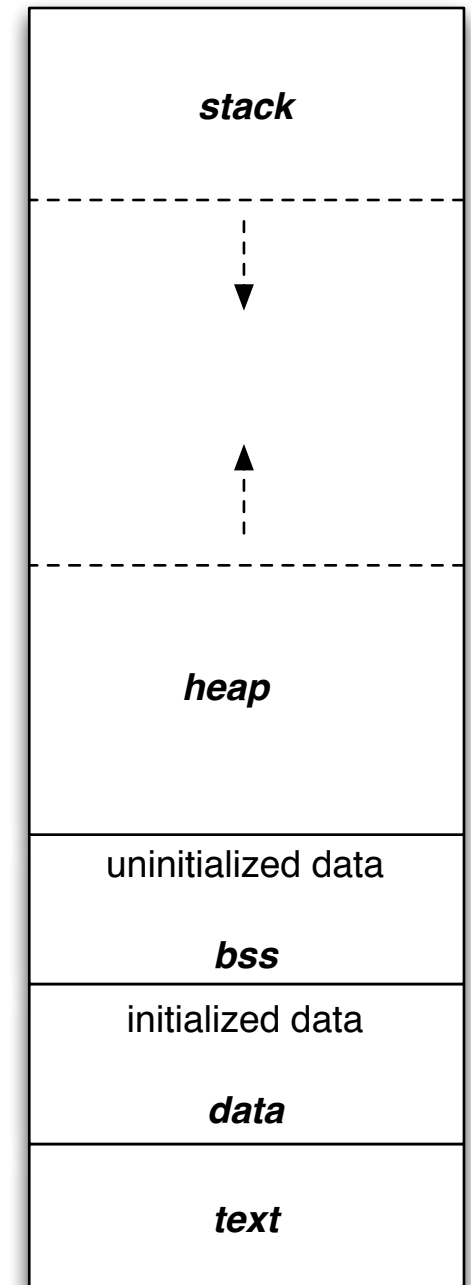
```
; nasm -f elf64 -F dwarf -g sbrk.asm && ld -o sbrk sbrk.o && gdb sbrk
```

```
section .data
initial_break: dq 0
current_break: dq 0
```

```
section .text
global _start
_start:
```

```
; get current break address
```

```
;-----
    mov     rax, 12                ;system call brk
    mov     rdi, 0                ;invalid address
    syscall
    mov     [current_break], rax
    mov     [initial_break], rax
```



Heap?

```

; allocate another 104857600 bytes of memory on the heap
;-----
    mov     rax, 12                ;system call brk
    mov     rdi, [current_break]
    add     rdi, 104857600         ;allocate 104857600 bytes
    syscall

.b0: ;Break the program here in GDB. Also watch the memory used by
    ;using "top" command. You'll noticed the memory used by this
    ;program dropped from 100M to 160k.
    mov     [current_break], rax


    sub     rax, 8
    mov     qword [rax], 123       ;write a qword to the end of heap

; free all allocated memory on the heap
;-----
    mov     rax, 12                ;system call brk
    mov     rdi, [initial_break] ;reset break address to its initial addr
    syscall

.b1: ;Break the program here in GDB, to see the memory drop.

.exit:
    mov     rax, 60 ;system call exit
    mov     rdi, 0 ;return value := 0
    syscall

```



Learn NASM by Disassembly

- c2nasm.bash: <https://gist.github.com/abcdabcd987/acb76b101094edac57537ab54ef1c4ef>
- `./c2nasm.bash program.c`
- `cat program.asm`

Summary

- Linux System Call
- C Library
- NASM Format
- x86-64 Registers, Memory Addressing, Instruction Set
- x86-64 SystemV Calling Convention